

AI Alone Will Not Modernize the Life Sciences Enterprise

AI can accelerate decisions, automate workflows, and uncover insights at scale. But without trusted data, semantic context, and governance, it can also accelerate confusion

ADAM SILVIS, Technology Strategy Lead, IQVIA

For life sciences executives, the hard truth is clear: technology spend does not equal enterprise modernization. Modernization depends on whether the organization can combine trusted data, scalable technology, and governance to enable decisions at enterprise scale. The companies that create durable value will not be those that run more pilots, build more dashboards, or deploy more AI tools. They will be the ones that make complex data understandable, trusted, governed, and actionable across the enterprise while giving teams the discipline and speed to act.

That distinction matters because the commercial mandate has not changed. Life sciences organizations still need to prepare more effectively for high-value customer interactions, allocate spend with greater precision, understand care pathways and patient journeys, and respond to access barriers and customer behavior. What has changed is the tolerance for delay and ambiguity. Senior decision makers need answers faster, with context they can trust and enough confidence to act.

Agentic AI and other advanced technologies raise the stakes because they can interpret, recommend, orchestrate, and trigger action across workflows. But they cannot compensate for weak enterprise foundations. If data is fragmented, definitions are inconsistent, quality is questionable, and governance is weak, AI will not transform the enterprise. It will expose the disconnects leaders have not resolved.

The use cases are obvious. The operating model is the constraint.

Across the commercial enterprise, the problem is not a lack of imagination. The use cases are clear: sharper field preparation, more responsive omnichannel planning, better spend allocation, improved customer understanding, faster insight generation, and more precise engagement. The constraint is execution. Too many initiatives remain trapped between promising pilots and scaled enterprise value because the operating model has not caught up to the technology agenda.


This is where C-suite attention matters. Models will evolve. Costs will change. Routing strategies will mature. New tools will enter the market. But the differentiating capability is the enterprise's ability to connect data, define it consistently, govern it transparently, and create decision processes that convert insight into action with confidence and accountability.

For many organizations, access to data is no longer the only bottleneck. The harder question is whether the data can be understood in context. A field name, metric, or KPI may be familiar to an experienced analyst, but an AI agent needs explicit meaning. If a commercial metric such as new-to-brand performance over a trailing time period is represented only by shorthand, the system must be able to connect that shorthand to business meaning, lineage, ownership, quality, and approved usage.

Interpretability is no longer a technical preference. It is an executive requirement. In an agentic environment, data must be more than available. It must be explainable, trusted, and usable in the flow of decision making.

Bad data is not a back-office problem. It is a commercial liability.

In life sciences, data quality has always mattered. In the agentic era, it has become a front-line commercial issue. Poor data quality is no longer simply a reporting inconvenience or remediation task. It can translate into the wrong customer being prioritized, the wrong message being delivered, the wrong channel being selected, or the wrong signal being interpreted as an opportunity.



AI changes the economics of data quality. Historically, poor-quality data often produced isolated reporting errors. AI can operationalize those errors across hundreds or thousands of decisions simultaneously. The question is no longer whether bad data creates risk. The question is how quickly that risk can scale.

As commercial teams move toward more personalized and dynamic engagement, the margin for ambiguity narrows. Omnichannel execution depends on the ability to connect touchpoints across field, digital, service, access, and engagement channels. If those touchpoints are not governed through a common data spine, teams may gain speed in isolated pockets while increasing fragmentation across the enterprise.

This is the paradox many organizations now face: the pursuit of speed and local optimization can create the fragmentation that prevents enterprise scale. Brand teams want faster answers. Field teams want more relevant recommendations. Analytics teams want to activate new signals. Technology teams want to enable experimentation. But when each function solves for speed independently, the enterprise can end up with siloed processes, inconsistent definitions, duplicated logic, and uneven trust.

The answer is not to slow innovation. It is to distinguish between what must be standardized and where teams can adapt and experiment. Core customer, product, territory, consent, compliance, and performance data require disciplined governance. Execution layers, including campaign design, content personalization, channel orchestration, and local experimentation, can then move faster because they are anchored to trusted foundations.

Governance cannot stop at cataloging assets. It must govern decisions.

Traditional governance disciplines remain essential, but they are no longer sufficient on their own. Data catalogs, glossaries, taxonomies, and metadata repositories are important building blocks. Yet, decision makers increasingly need visibility across the full lifecycle of data: where it originated, how it was transformed, which business rules were applied, what quality checks were performed, who owns it, where it is used, and what compliance exposure it may carry.

In traditional analytics environments, governance primarily supported reporting and compliance. In AI environments, governance becomes part of the reasoning process itself. Agents must be able to understand not only what data exists, but whether that data is trusted, current, authorized, and appropriate for the decision being made.

The executive question is not simply, “Do we have this data?” It is, “Can I trust the answer this data produced?” If a senior leader asks why a key performance indicator changed in a specific region, the organization should

not need to assemble a cross-functional team for a week simply to trace the answer. The goal is not another dashboard. The goal is decision-ready data that can be queried, interpreted, explained, and acted on with confidence.

This requires governance to become more connected, more semantic, and more operational. Ontologies, knowledge graphs, metadata lineage, business context, and temporal history all become more important when AI agents are expected to reason across data assets. A current master data record may show the latest state of a customer, but many business questions require understanding how that customer changed over time. Without temporal context, semantic clarity, and end-to-end lineage, agents may have access to information but lack the context required to use it responsibly.

For regulated life sciences organizations, this is not optional. Data assets used to support automation, analytics, and AI must be trusted, auditable, and aligned to the appropriate business and compliance controls. Analytics-ready data cannot become ungoverned data simply because the enterprise wants to move faster.



If technology does not fit the workflow, it will not change the business.

A compelling pilot is not the same as an enterprise capability. Many technology initiatives stall because they are impressive in demonstration but disconnected from how work actually gets done. If users must leave their normal workflow, go to a separate destination, interpret a separate output, and manually translate that insight into action, the enterprise has added friction rather than removed it.

The most durable use cases are embedded where decisions already happen: in customer relationship management systems, omnichannel planning tools, field preparation workflows, analytics environments, and mobile experiences that support customer-facing teams. Technology must reduce friction, not add another interface to manage.

Leadership behavior also matters. Enterprise transformation does not occur when senior leaders simply instruct teams to “use AI”. It occurs when leaders model responsible experimentation, acknowledge uncertainty, ask better questions, and create permission for teams to learn while maintaining clear standards. In organizations where leaders actively engage with technology, data quality, governance, and responsible AI use, adoption becomes less performative and more practical.

Distributed creation without accountability becomes enterprise risk.

As AI-enabled development tools become more accessible, the definition of “builder” is expanding. Employees who never considered themselves developers can now create applications, automations, analyses, and workflows through natural language and assisted coding tools. That creates meaningful opportunity: more innovation, faster prototyping, and closer alignment between business needs and technical output.

It also creates a new governance challenge. When creation scales from a controlled group of developers to a much broader population of business users, organizations need new frameworks for ownership, review, maintenance, risk assessment, and lifecycle management. A tool built quickly for one team may affect downstream processes, data quality, compliance posture, or customer engagement if it is connected to shared systems.

Business users should be empowered. But empowerment cannot outpace accountability. Modern enterprise enablement requires clear human ownership, defined boundaries, transparent auditability, and a governance model that can keep pace with distributed innovation.



Scaled value requires executive choices, not more experimentation.

For life sciences decision makers, the path forward requires discipline. Organizations will not scale value by continuing to multiply isolated pilots. They will scale value by making explicit executive choices about what must be standardized, where experimentation should be encouraged, and how data, AI, and technology will be governed in the flow of business.



First, protect the foundation.

Standardize and govern the core data assets elements that define customers, products, territories, access, consent, performance, and engagement. These are not areas for uncontrolled variation. They are the connective tissue of commercial execution.



Second, make data explainable.

Build semantic layers that connect technical fields to business meaning. Ensure that agents and users can understand definitions, lineage, quality, approved uses, and the basis for AI-generated outputs without relying on tribal knowledge.



Third, govern decisions, not just data.

Move beyond cataloging assets toward governing the decisions, recommendations, automated actions, and workflows those assets support. Trust must be visible at the point of use.



Fourth, embed trust into real workflows. Prioritize use cases that meet users where they already work and make data quality, lineage, and approved usage visible at the point of decision. Adoption improves when technology makes existing tasks faster, clearer, or more effective without asking users to validate every output independently.



Fifth, scale responsibly.

Encourage experimentation at the edge while maintaining disciplined governance over shared data, approved logic, compliance controls, production workflows, and AI-enabled actions. Innovation should accelerate enterprise priorities, not create another layer of fragmentation.

The prize is not modernization. It is decision advantage.

The most compelling promise of modern technology in life sciences is not simply doing the same work with fewer steps. It is enabling organizations to ask better questions, connect more signals, reduce decision latency, and act with greater confidence across the value chain.

That promise will not be realized through technology enthusiasm alone. It will require data foundations that are connected and governed, semantic context that makes information usable, workflows that bring intelligence to the point of decision, and governance models that ensure AI operates on trusted, explainable, and auditable information.

Life sciences leaders should resist the temptation to treat technology as a shortcut around foundational work. AI will not create trust where trust does not already exist. It will not create context where context is missing. And it will not create quality where quality has not been designed into the enterprise. The foundation is the strategy. Organizations that invest in data quality, governance, and semantic understanding today will be best positioned to scale AI into a true decision advantage tomorrow.